

HR Attrition Analysis Using Random Forest

AhmedLubis

2026-07-04

HUMAN RESOURCES ATTRITION ANALYSIS USING RANDOM FOREST

Step 1: Load Required Libraries

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.5.3
```

```
## Warning: package 'ggplot2' was built under R version 4.5.3
```

```
## Warning: package 'purrr' was built under R version 4.5.3
```

```
## Warning: package 'dplyr' was built under R version 4.5.3
```

```
## — Attaching core tidyverse packages — tidyverse 2.0.0 —
## ✓ dplyr      1.2.1    ✓ readr      2.1.6
## ✓ forcats    1.0.1    ✓ stringr    1.6.0
## ✓ ggplot2    4.0.3    ✓ tibble     3.3.1
## ✓ lubridate  1.9.5    ✓ tidyr      1.3.2
## ✓ purrr      1.2.2
## — Conflicts — tidyverse_conflicts() —
## X dplyr::filter() masks stats::filter()
## X dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 4.5.3
```

```
## randomForest 4.7-1.2
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##   combine
##
## The following object is masked from 'package:ggplot2':
##
##   margin
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.5.3
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift
```

Step 2: Load and Clean the Dataset

```
# Read the file with proper delimiter (;) and handle empty text inputs as NA
hr_raw <- read_delim("hr_data.csv", delim = ";", na = c("", "NA"))
```

```
## Rows: 14999 Columns: 10
## — Column specification —————
## Delimiter: ";"
## chr (2): sales, salary
## dbl (8): satisfaction_level, last_evaluation, number_project, average_monthly...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
# Clean columns and handle categorical data types
hr_cleaned <- hr_raw %>%
  rename(department = sales) %>%
  drop_na() %>% # Drop rows with missing values to ensure model stability
  mutate(
    # Target variable 'left' must be a factor for classification
    left = factor(left, levels = c(0, 1), labels = c("Stayed", "Left")),

    # Format all categorical variables as factors
    salary = factor(salary, levels = c("low", "medium", "high")),
    department = as.factor(department),
    Work_accident = as.factor(Work_accident),
    promotion_last_5years = as.factor(promotion_last_5years)
  )

# Verify structural transformations
glimpse(hr_cleaned)
```

```
## Rows: 14,430
## Columns: 10
## $ satisfaction_level    <dbl> 0.80, 0.72, 0.37, 0.41, 0.10, 0.92, 0.89, 0.42, ...
## $ last_evaluation       <dbl> 0.86, 0.87, 0.52, 0.50, 0.77, 0.85, 1.00, 0.53, ...
## $ number_project        <dbl> 5, 5, 2, 2, 6, 5, 5, 2, 2, 6, 4, 2, 2, 2, 4, ...
## $ average_monthly_hours <dbl> 262, 223, 159, 153, 247, 259, 224, 142, 135, 305...
## $ time_spend_company    <dbl> 6, 5, 3, 3, 4, 5, 5, 3, 3, 4, 5, 3, 3, 3, 3, 6, ...
## $ Work_accident         <fct> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ left                  <fct> Left, Left, Left, Left, Left, Left, Left, Left, ...
## $ promotion_last_5years <fct> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ department            <fct> sales, sales, sales, sales, sales, sales, sales, ...
## $ salary                <fct> medium, low, low, low, low, low, low, low, low, ...
```

Step 3: Split Data into Training and Testing Sets

```
set.seed(123)

# Stratified split: 70% for training, 30% for testing
train_index <- createDataPartition(hr_cleaned$left, p = 0.70, list = FALSE)
train_set   <- hr_cleaned[train_index, ]
test_set    <- hr_cleaned[-train_index, ]

# Print sizes to confirm the split balance
cat("Training Set size:", nrow(train_set), "\n")
```

```
## Training Set size: 10102
```

```
cat("Testing Set size:", nrow(test_set), "\n")
```

```
## Testing Set size: 4328
```

Step 4: Train the Random Forest Classifier

```
# Train an ensemble of 500 decision trees
rf_model <- randomForest(
  left ~ .,
  data = train_set,
  ntree = 500,
  importance = TRUE
)

# Print overall model summary and Out-Of-Bag (OOB) error rate
print(rf_model)
```

```
##
## Call:
## randomForest(formula = left ~ ., data = train_set, ntree = 500,      importance = TRUE)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 3
##
##              OOB estimate of  error rate: 1%
## Confusion matrix:
##              Stayed Left class.error
## Stayed   7987   13  0.00162500
## Left      88 2014  0.04186489
```

Step 5: Evaluate Model Performance (Confusion Matrix)

```
# Predict the status of employees in the unseen test set
test_predictions <- predict(rf_model, newdata = test_set)

# Generate a comprehensive confusion matrix report
conf_matrix <- confusionMatrix(test_predictions, test_set$left)
print(conf_matrix)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Stayed Left
##   Stayed   3422   48
##   Left      6   852
##
##           Accuracy : 0.9875
##           95% CI : (0.9838, 0.9906)
##   No Information Rate : 0.7921
##   P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9615
##
##   Mcnemar's Test P-Value : 2.414e-08
##
##           Sensitivity : 0.9982
##           Specificity : 0.9467
##   Pos Pred Value : 0.9862
##   Neg Pred Value : 0.9930
##   Prevalence : 0.7921
##   Detection Rate : 0.7907
##   Detection Prevalence : 0.8018
##   Balanced Accuracy : 0.9725
##
##   'Positive' Class : Stayed
##
```

Step 6: Extract and Plot Feature Importance

```
# Display the numerical importance scores for each variable
print(importance(rf_model))
```

```
##           Stayed      Left MeanDecreaseAccuracy MeanDecreaseGini
## satisfaction_level  75.675891 218.05560          215.97178      1124.478170
## last_evaluation    22.238364 134.10255          134.67260       419.977021
## number_project     44.378719 144.03300          142.18687       581.122400
## average_monthly_hours 60.465122 88.37713           96.37217       484.799451
## time_spend_company  62.978753 86.62907           95.23014       593.849100
## Work_accident       8.733582 22.19413           21.76427        19.823221
## promotion_last_5years 6.818380 13.91332           14.65314         3.153431
## department         12.983583 55.16677           37.62590        56.310925
## salary              15.097616 37.68420           33.98818        29.594204
```

```
# Generate visual charts for Mean Decrease Accuracy and Mean Decrease Gini
varImpPlot(rf_model, main = "Feature Importance for Employee Attrition")
```

Feature Importance for Employee Attrition

